

In this tutorial consider the CIFAR dataset. Use the initial code given below and perform the classification using CNN and HoG based features.

```
In [ ]: # Auto-setup when running on Google Colab
if 'google.colab' in str(get_ipython()):
    !pip install openml

# General imports
%matplotlib inline
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import openml as oml
import tensorflow as tf
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: openml in /usr/local/lib/python3.10/dist-packages (0.13.1)
Requirement already satisfied: liac-arff>=2.4.0 in /usr/local/lib/python3.10/dist-packages (from openml) (2.5.0)
Requirement already satisfied: xlrd in /usr/local/lib/python3.10/dist-packages (from openml) (2.0.1)
Requirement already satisfied: xlsxwriter in /usr/local/lib/python3.10/dist-packages (from openml) (3.0.4)
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from openml) (2.27.1)
Requirement already satisfied: scikit-learn>=0.18 in /usr/local/lib/python3.10/dist-packages (from openml) (1.2.2)
Requirement already satisfied: python-dateutil in /usr/local/lib/python3.10/dist-packages (from openml) (2.8.2)
Requirement already satisfied: pandas>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from openml) (1.5.3)
Requirement already satisfied: scipy>=0.13.3 in /usr/local/lib/python3.10/dist-packages (from openml) (1.10.1)
Requirement already satisfied: numpy>=1.6.2 in /usr/local/lib/python3.10/dist-packages (from openml) (1.22.4)
Requirement already satisfied: minio in /usr/local/lib/python3.10/dist-packages (from openml) (7.1.15)
Requirement already satisfied: pyarrow in /usr/local/lib/python3.10/dist-packages (from openml) (9.0.0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=1.0.0->openml) (2022.7.1)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil->openml) (1.16.0)
Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.18->openml) (1.2.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.18->openml) (3.1.0)
Requirement already satisfied: certifi in /usr/local/lib/python3.10/dist-packages (from minio->openml) (2022.12.7)
Requirement already satisfied: urllib3 in /usr/local/lib/python3.10/dist-packages (from minio->openml) (1.26.15)
Requirement already satisfied: charset-normalizer<=2.0.0 in /usr/local/lib/python3.10/dist-packages (from requests->openml) (2.0.12)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests->openml) (3.4)
```

```
In [ ]: # Uncomment the next line if you run on Colab
#!pip install --quiet openml
```

```
In [ ]: %matplotlib inline
import openml as oml
import matplotlib.pyplot as plt
```

```
In [ ]: # Download CIFAR data. Takes a while the first time.
# This version returns 3x32x32 resolution images.
# If you feel like it, repeat the exercises with the 96x96x3 resolution version by using ID 41103
cifar = oml.datasets.get_dataset(40926)
X, y, _, _ = cifar.get_data(target=cifar.default_target_attribute, dataset_format='array');
cifar_classes = {0: "airplane", 1: "automobile", 2: "bird", 3: "cat", 4: "deer",
                 5: "dog", 6: "frog", 7: "horse", 8: "ship", 9: "truck"}
```

```
In [ ]: # The dataset (40926) is in a weird 3x32x32 format, we need to reshape and transpose
Xr = X.reshape((len(X),3,32,32)).transpose(0,2,3,1)
```

```
In [ ]: Xr
```

```
Out[ ]: array([[[[ 59.,  62.,  63.],
               [ 43.,  46.,  45.],
               [ 50.,  48.,  43.],
               ...,
               [158., 132., 108.],
               [152., 125., 102.],
               [148., 124., 103.]],

               [[ 16.,  20.,  20.],
               [  0.,   0.,   0.],
               [ 18.,   8.,   0.],
               ...,
               [123.,  88.,  55.],
               [119.,  83.,  50.],
               [122.,  87.,  57.]],

               [[ 25.,  24.,  21.],
               [ 16.,   7.,   0.]])])
```

```

[ 49., 27., 8.],
...
[118., 84., 50.],
[120., 84., 50.],
[109., 73., 42.]],

...,

[[208., 170., 96.],
 [201., 153., 34.],
 [198., 161., 26.],
 ...
 [160., 133., 70.],
 [ 56., 31., 7.],
 [ 53., 34., 20.]],

[[180., 139., 96.],
 [173., 123., 42.],
 [186., 144., 30.],
 ...
 [184., 148., 94.],
 [ 97., 62., 34.],
 [ 83., 53., 34.]],

[[177., 144., 116.],
 [168., 129., 94.],
 [179., 142., 87.],
 ...
 [216., 184., 140.],
 [151., 118., 84.],
 [123., 92., 72.]]],

[[[154., 177., 187.],
  [126., 137., 136.],
  [105., 104., 95.],
  ...
  [ 91., 95., 71.],
  [ 87., 90., 71.],
  [ 79., 81., 70.]],

[[140., 160., 169.],
 [145., 153., 154.],
 [125., 125., 118.],
 ...
 [ 96., 99., 78.],
 [ 77., 80., 62.],
 [ 71., 73., 61.]],

[[140., 155., 164.],
 [139., 146., 149.],
 [115., 115., 112.],
 ...
 [ 79., 82., 64.],
 [ 68., 70., 55.],
 [ 67., 69., 55.]],

...,

[[175., 167., 166.],
 [156., 154., 160.],
 [154., 160., 170.],
 ...
 [ 42., 34., 36.],
 [ 61., 53., 57.],
 [ 93., 83., 91.]],

[[165., 154., 128.],
 [156., 152., 130.],
 [159., 161., 142.],
 ...
 [103., 93., 96.],
 [123., 114., 120.],
 [131., 121., 131.]],

[[163., 148., 120.],
 [158., 148., 122.],
 [163., 156., 133.],
 ...
 [143., 133., 139.],
 [143., 134., 142.],
 [143., 133., 144.]]],

[[[255., 255., 255.],
  [253., 253., 253.],
  [253., 253., 253.],
  ...
  [253., 253., 253.],

```

```

[253., 253., 253.],
[253., 253., 253.]],

[[255., 255., 255.],
[255., 255., 255.],
[255., 255., 255.],
...,
[255., 255., 255.],
[255., 255., 255.],
[255., 255., 255.]],

[[255., 255., 255.],
[254., 254., 254.],
[254., 254., 254.],
...,
[254., 254., 254.],
[254., 254., 254.],
[254., 254., 254.]],

...,

[[113., 120., 112.],
[111., 118., 111.],
[105., 112., 106.],
...,
[ 72.,  81.,  80.],
[ 72.,  80.,  79.],
[ 72.,  80.,  79.]],

[[111., 118., 110.],
[104., 111., 104.],
[ 99., 106.,  98.],
...,
[ 68.,  75.,  73.],
[ 70.,  76.,  75.],
[ 78.,  84.,  82.]],

[[106., 113., 105.],
[ 99., 106.,  98.],
[ 95., 102.,  94.],
...,
[ 78.,  85.,  83.],
[ 79.,  85.,  83.],
[ 80.,  86.,  84.]]],

...,

[[[ 20.,  15.,  12.],
[ 19.,  14.,  11.],
[ 15.,  14.,  11.],
...,
[ 10.,   9.,   7.],
[ 12.,  11.,   9.],
[ 13.,  12.,  10.]],

[[ 21.,  16.,  13.],
[ 20.,  16.,  13.],
[ 18.,  17.,  12.],
...,
[ 10.,   9.,   7.],
[ 10.,   9.,   7.],
[ 12.,  11.,   9.]],

[[ 21.,  16.,  13.],
[ 21.,  17.,  12.],
[ 20.,  18.,  11.],
...,
[ 12.,  11.,   9.],
[ 12.,  11.,   9.],
[ 13.,  12.,  10.]],

...,

[[ 33.,  25.,  13.],
[ 34.,  26.,  15.],
[ 34.,  26.,  15.],
...,
[ 28.,  25.,  52.],
[ 29.,  25.,  58.],
[ 23.,  20.,  42.]],

[[ 33.,  25.,  14.],
[ 34.,  26.,  15.],
[ 34.,  26.,  15.],
...,
[ 27.,  24.,  52.],
[ 27.,  24.,  56.],

```

```

[ 25., 22., 47.]],

[[ 31., 23., 12.],
 [ 32., 24., 13.],
 [ 33., 25., 14.],
 ...,
 [ 24., 23., 50.],
 [ 26., 23., 53.],
 [ 25., 20., 47.]]],

[[[ 25., 40., 12.],
 [ 15., 36., 3.],
 [ 23., 41., 18.],
 ...,
 [ 61., 82., 78.],
 [ 92., 113., 112.],
 [ 75., 89., 92.]],

[[ 12., 25., 6.],
 [ 20., 37., 7.],
 [ 24., 36., 15.],
 ...,
 [115., 134., 138.],
 [149., 168., 177.],
 [104., 117., 131.]],

[[ 12., 25., 11.],
 [ 15., 29., 6.],
 [ 34., 40., 24.],
 ...,
 [154., 172., 182.],
 [157., 175., 192.],
 [116., 129., 151.]],

...,

[[100., 129., 81.],
 [103., 132., 84.],
 [104., 134., 86.],
 ...,
 [ 97., 128., 84.],
 [ 98., 126., 84.],
 [ 91., 121., 79.]],

[[103., 132., 83.],
 [104., 131., 83.],
 [107., 135., 87.],
 ...,
 [101., 132., 87.],
 [ 99., 127., 84.],
 [ 92., 121., 79.]],

[[ 95., 126., 78.],
 [ 95., 123., 76.],
 [101., 128., 81.],
 ...,
 [ 93., 124., 80.],
 [ 95., 123., 81.],
 [ 92., 120., 80.]]],

[[[ 73., 78., 75.],
 [ 98., 103., 113.],
 [ 99., 106., 114.],
 ...,
 [135., 150., 152.],
 [135., 149., 154.],
 [203., 215., 223.]],

[[ 69., 73., 70.],
 [ 84., 89., 97.],
 [ 68., 75., 81.],
 ...,
 [ 85., 95., 89.],
 [ 71., 82., 80.],
 [120., 133., 135.]],

[[ 69., 73., 70.],
 [ 90., 95., 100.],
 [ 62., 71., 74.],
 ...,
 [ 74., 81., 70.],
 [ 53., 62., 54.],
 [ 62., 74., 69.]],

...,

[[123., 128., 96.],

```

```

[132., 132., 102.],
[129., 128., 100.],
...,
[108., 107., 88.],
[ 62., 60., 55.],
[ 27., 27., 28.]],

[[115., 121., 91.],
[123., 124., 95.],
[129., 126., 99.],
...,
[115., 116., 94.],
[ 66., 65., 59.],
[ 27., 27., 27.]],


[[116., 120., 90.],
[121., 122., 94.],
[129., 128., 101.],
...,
[116., 115., 94.],
[ 68., 65., 58.],
[ 27., 26., 26.]]], dtype=float32)

```

```

In [ ]: # Take some random examples, reshape to a 32x32 image and plot
from random import randint
fig, axes = plt.subplots(1, 5, figsize=(10, 5))
for i in range(5):
    n = randint(0, len(Xr))
    # The data is stored in a 3x32x32 format, so we need to transpose it
    axes[i].imshow(Xr[n]/255)
    axes[i].set_xlabel((cifar_classes[int(y[n])]))
    axes[i].set_xticks((), axes[i].set_yticks(())
plt.show();

```



horse




truck



horse



automobile



deer

## Task 1: CNN based Classification

- Split the data into 80% training and 20% validation sets
- Normalize the data to [0,1]
- Build a ConvNet with 3 convolutional layers interspersed with MaxPooling layers, and one dense layer.
  - Use at least 32 3x3 filters in the first layer and ReLU activation.
  - Otherwise, make rational design choices or experiment a bit to see what works.
- You should at least get 60% accuracy.
- For training, you can try 20-50 epochs, but feel free to explore this as well

### Splitting the data into 80% training and 20% validation

```

In [ ]: from sklearn.model_selection import train_test_split

# Split the data into 80% training and 20% validation sets
X_train, X_val, y_train, y_val = train_test_split(Xr, y, test_size=0.2, random_state=42)
X_train.shape

```

```

Out[ ]: (16000, 32, 32, 3)

```

```

In [ ]: y_train.shape

```

```

Out[ ]: (16000,)

```

```

In [ ]: #normalizing the data
X_train = X_train / 255.0
X_val = X_val / 255.0

```

```

In [ ]: # Build the ConvNet model
model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(16, (3, 3), activation='relu', input_shape=(32, 32, 3)),
    tf.keras.layers.MaxPooling2D((2, 2)),
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D((2, 2)),
    tf.keras.layers.Conv2D(128, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D((2, 2)),
    tf.keras.layers.Flatten(),

```

```
tf.keras.layers.Dense(256, activation='relu'),
tf.keras.layers.Dense(10, activation='softmax')
])
```

```
In [ ]: # Compile the model
model.compile(optimizer='adam',
              loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
              metrics=['accuracy'])
```

```
In [ ]: # Train the model
history = model.fit(X_train, y_train, epochs=50, validation_data=(X_val, y_val))

# Evaluate the model on the validation set
val_loss, val_acc = model.evaluate(X_val, y_val)
```

Epoch 1/50

/usr/local/lib/python3.10/dist-packages/keras/backend.py:5612: UserWarning: "`sparse\_categorical\_crossentropy` received `from\_logits=True`, but the `output` argument was produced by a Softmax activation and thus does not represent logits. Was this intended?

output, from\_logits = \_get\_logits(

500/500 [=====] - 5s 5ms/step - loss: 1.7534 - accuracy: 0.3464 - val\_loss: 1.5242 - val\_accuracy: 0.4540

Epoch 2/50

500/500 [=====] - 2s 4ms/step - loss: 1.3907 - accuracy: 0.4971 - val\_loss: 1.4140 - val\_accuracy: 0.4927

Epoch 3/50

500/500 [=====] - 2s 4ms/step - loss: 1.2357 - accuracy: 0.5603 - val\_loss: 1.3377 - val\_accuracy: 0.5185

Epoch 4/50

500/500 [=====] - 2s 4ms/step - loss: 1.1289 - accuracy: 0.5969 - val\_loss: 1.2053 - val\_accuracy: 0.5720

Epoch 5/50

500/500 [=====] - 3s 7ms/step - loss: 1.0311 - accuracy: 0.6363 - val\_loss: 1.2326 - val\_accuracy: 0.5690

Epoch 6/50

500/500 [=====] - 2s 4ms/step - loss: 0.9418 - accuracy: 0.6641 - val\_loss: 1.1241 - val\_accuracy: 0.6093

Epoch 7/50

500/500 [=====] - 2s 4ms/step - loss: 0.8618 - accuracy: 0.6958 - val\_loss: 1.1473 - val\_accuracy: 0.6090

Epoch 8/50

500/500 [=====] - 2s 4ms/step - loss: 0.7907 - accuracy: 0.7179 - val\_loss: 1.1277 - val\_accuracy: 0.6295

Epoch 9/50

500/500 [=====] - 2s 4ms/step - loss: 0.7125 - accuracy: 0.7470 - val\_loss: 1.1831 - val\_accuracy: 0.6175

Epoch 10/50

500/500 [=====] - 3s 5ms/step - loss: 0.6410 - accuracy: 0.7731 - val\_loss: 1.1764 - val\_accuracy: 0.6320

Epoch 11/50

500/500 [=====] - 2s 5ms/step - loss: 0.5587 - accuracy: 0.8075 - val\_loss: 1.2536 - val\_accuracy: 0.6233

Epoch 12/50

500/500 [=====] - 2s 4ms/step - loss: 0.4978 - accuracy: 0.8253 - val\_loss: 1.3208 - val\_accuracy: 0.6165

Epoch 13/50

500/500 [=====] - 2s 4ms/step - loss: 0.4419 - accuracy: 0.8464 - val\_loss: 1.3979 - val\_accuracy: 0.6168

Epoch 14/50

500/500 [=====] - 2s 4ms/step - loss: 0.3806 - accuracy: 0.8686 - val\_loss: 1.4438 - val\_accuracy: 0.6235

Epoch 15/50

500/500 [=====] - 3s 5ms/step - loss: 0.3323 - accuracy: 0.8823 - val\_loss: 1.5844 - val\_accuracy: 0.5985

Epoch 16/50

500/500 [=====] - 3s 5ms/step - loss: 0.2744 - accuracy: 0.9057 - val\_loss: 1.6014 - val\_accuracy: 0.6227

Epoch 17/50

500/500 [=====] - 2s 5ms/step - loss: 0.2325 - accuracy: 0.9209 - val\_loss: 1.8256 - val\_accuracy: 0.6045

Epoch 18/50

500/500 [=====] - 2s 4ms/step - loss: 0.2132 - accuracy: 0.9256 - val\_loss: 1.8051 - val\_accuracy: 0.6190

Epoch 19/50

500/500 [=====] - 4s 8ms/step - loss: 0.1715 - accuracy: 0.9397 - val\_loss: 2.0009 - val\_accuracy: 0.6145

Epoch 20/50

500/500 [=====] - 5s 10ms/step - loss: 0.1422 - accuracy: 0.9534 - val\_loss: 2.0633 - val\_accuracy: 0.6160

Epoch 21/50

500/500 [=====] - 3s 7ms/step - loss: 0.1336 - accuracy: 0.9556 - val\_loss: 2.1868 - val\_accuracy: 0.6223

Epoch 22/50

500/500 [=====] - 2s 4ms/step - loss: 0.1338 - accuracy: 0.9549 - val\_loss: 2.1824 - val\_accuracy: 0.6090

Epoch 23/50

500/500 [=====] - 2s 5ms/step - loss: 0.1179 - accuracy: 0.9588 - val\_loss: 2.3605 - val\_accuracy: 0.6165

```

Epoch 24/50
500/500 [=====] - 3s 5ms/step - loss: 0.1009 - accuracy: 0.9664 - val_loss: 2.5011 - v
al_accuracy: 0.6047
Epoch 25/50
500/500 [=====] - 2s 5ms/step - loss: 0.1010 - accuracy: 0.9653 - val_loss: 2.5552 - v
al_accuracy: 0.6110
Epoch 26/50
500/500 [=====] - 2s 4ms/step - loss: 0.0906 - accuracy: 0.9692 - val_loss: 2.6882 - v
al_accuracy: 0.6137
Epoch 27/50
500/500 [=====] - 2s 4ms/step - loss: 0.1114 - accuracy: 0.9632 - val_loss: 2.8709 - v
al_accuracy: 0.5993
Epoch 28/50
500/500 [=====] - 3s 6ms/step - loss: 0.0921 - accuracy: 0.9684 - val_loss: 2.6450 - v
al_accuracy: 0.6105
Epoch 29/50
500/500 [=====] - 3s 6ms/step - loss: 0.0697 - accuracy: 0.9774 - val_loss: 2.8983 - v
al_accuracy: 0.6008
Epoch 30/50
500/500 [=====] - 2s 5ms/step - loss: 0.0861 - accuracy: 0.9711 - val_loss: 3.1602 - v
al_accuracy: 0.5900
Epoch 31/50
500/500 [=====] - 2s 5ms/step - loss: 0.0769 - accuracy: 0.9735 - val_loss: 2.9133 - v
al_accuracy: 0.5985
Epoch 32/50
500/500 [=====] - 2s 4ms/step - loss: 0.0951 - accuracy: 0.9672 - val_loss: 3.0931 - v
al_accuracy: 0.6148
Epoch 33/50
500/500 [=====] - 2s 4ms/step - loss: 0.0701 - accuracy: 0.9753 - val_loss: 3.1975 - v
al_accuracy: 0.6115
Epoch 34/50
500/500 [=====] - 3s 6ms/step - loss: 0.0942 - accuracy: 0.9678 - val_loss: 3.0016 - v
al_accuracy: 0.6053
Epoch 35/50
500/500 [=====] - 2s 5ms/step - loss: 0.0535 - accuracy: 0.9827 - val_loss: 3.1260 - v
al_accuracy: 0.6045
Epoch 36/50
500/500 [=====] - 2s 5ms/step - loss: 0.0596 - accuracy: 0.9803 - val_loss: 3.1107 - v
al_accuracy: 0.6120
Epoch 37/50
500/500 [=====] - 2s 4ms/step - loss: 0.0665 - accuracy: 0.9771 - val_loss: 3.4090 - v
al_accuracy: 0.6058
Epoch 38/50
500/500 [=====] - 2s 4ms/step - loss: 0.0743 - accuracy: 0.9743 - val_loss: 3.2444 - v
al_accuracy: 0.6170
Epoch 39/50
500/500 [=====] - 3s 5ms/step - loss: 0.0607 - accuracy: 0.9804 - val_loss: 3.3561 - v
al_accuracy: 0.5950
Epoch 40/50
500/500 [=====] - 3s 5ms/step - loss: 0.0804 - accuracy: 0.9728 - val_loss: 3.3980 - v
al_accuracy: 0.5925
Epoch 41/50
500/500 [=====] - 2s 5ms/step - loss: 0.0587 - accuracy: 0.9808 - val_loss: 3.3581 - v
al_accuracy: 0.6227
Epoch 42/50
500/500 [=====] - 2s 5ms/step - loss: 0.0706 - accuracy: 0.9772 - val_loss: 3.4136 - v
al_accuracy: 0.6053
Epoch 43/50
500/500 [=====] - 2s 4ms/step - loss: 0.0682 - accuracy: 0.9769 - val_loss: 3.4453 - v
al_accuracy: 0.6083
Epoch 44/50
500/500 [=====] - 3s 5ms/step - loss: 0.0420 - accuracy: 0.9859 - val_loss: 3.6554 - v
al_accuracy: 0.6155
Epoch 45/50
500/500 [=====] - 3s 5ms/step - loss: 0.0536 - accuracy: 0.9815 - val_loss: 4.1795 - v
al_accuracy: 0.5838
Epoch 46/50
500/500 [=====] - 2s 4ms/step - loss: 0.0981 - accuracy: 0.9671 - val_loss: 3.5745 - v
al_accuracy: 0.6028
Epoch 47/50
500/500 [=====] - 2s 5ms/step - loss: 0.0578 - accuracy: 0.9805 - val_loss: 3.5718 - v
al_accuracy: 0.6080
Epoch 48/50
500/500 [=====] - 2s 5ms/step - loss: 0.0305 - accuracy: 0.9899 - val_loss: 3.7301 - v
al_accuracy: 0.6065
Epoch 49/50
500/500 [=====] - 2s 5ms/step - loss: 0.0589 - accuracy: 0.9808 - val_loss: 3.7318 - v
al_accuracy: 0.6093
Epoch 50/50
500/500 [=====] - 3s 6ms/step - loss: 0.0852 - accuracy: 0.9722 - val_loss: 3.6693 - v
al_accuracy: 0.6012
125/125 [=====] - 0s 2ms/step - loss: 3.6693 - accuracy: 0.6012

```

```

In [ ]: ...
# list all data in history
print(history.history.keys())

dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])



```

```

In [ ]: print(history.history.keys())
# summarize history for accuracy
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
# summarize history for loss
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()

```

```
dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```



```

In [ ]: print("Validation Loss:", val_loss)
print("Validation Accuracy:", val_acc*100)

```

```

Validation Loss: 3.669343948364258
Validation Accuracy: 60.12499928474426

```

```

In [ ]: # Get the model's predictions on the validation set
y_val_pred = model.predict(X_val)
y_val_pred_labels = np.argmax(y_val_pred, axis=1)

# Calculate the accuracy by comparing the predicted labels with the true labels
val_accuracy = np.mean(y_val_pred_labels == y_val) * 100

```



```
# Print the validation accuracy
print("Validation Accuracy:", val_accuracy)
```

```
125/125 [=====] - 0s 2ms/step
Validation Accuracy: 48.55
```

## Task 2: HOG based Classification

- Use the same data split as mentioned above and extract the HOG features from the training and testing split and classify the HOG features using SVM classifier.
- Compare the accuracy with CNN features and explain which technique outperforms.

```
In [ ]: from sklearn.model_selection import train_test_split

# Split the data into 80% training and 20% validation sets
X_train, X_val, y_train, y_val = train_test_split(Xr, y, test_size=0.2, random_state=42)
```

```
In [ ]: from skimage.feature import hog

# Extract HOG features from the training set
X_train_hog = []
for image in X_train:
    features = hog(image, orientations=9, pixels_per_cell=(8, 8), cells_per_block=(2, 2), multichannel=True)
    X_train_hog.append(features)
X_train_hog = np.array(X_train_hog)
```

```
<ipython-input-54-91c041c72d67>:6: FutureWarning: `multichannel` is a deprecated argument name for `hog`. It will be removed in version 1.0. Please use `channel_axis` instead.
    features = hog(image, orientations=9, pixels_per_cell=(8, 8), cells_per_block=(2, 2), multichannel=True)
```

```
In [ ]: # Extract HOG features from the validation set
X_val_hog = []
for image in X_val:
    features = hog(image, orientations=9, pixels_per_cell=(8, 8), cells_per_block=(2, 2), multichannel=True)
    X_val_hog.append(features)
X_val_hog = np.array(X_val_hog)
```

```
<ipython-input-55-53234980e333>:4: FutureWarning: `multichannel` is a deprecated argument name for `hog`. It will be removed in version 1.0. Please use `channel_axis` instead.
    features = hog(image, orientations=9, pixels_per_cell=(8, 8), cells_per_block=(2, 2), multichannel=True)
```

```
In [ ]: from sklearn.svm import SVC

# Initialize and train the SVM classifier
svm_classifier = SVC()
svm_classifier.fit(X_train_hog, y_train)
accuracy_svm = svm_classifier.score(X_val_hog, y_val)
print("SVM Accuracy:", accuracy_svm)
```

```
SVM Accuracy: 0.5995
```


```
In [ ]: results=pd.DataFrame({'Model':['CNN','SVM'],
                             'Accuracy Score':[accuracy_cnn,accuracy_svm]})
result_df=results.sort_values(by='Accuracy Score', ascending=False)
result_df=result_df.set_index('Model')
result_df
```

```
Out[ ]:      Accuracy Score
Model
CNN          0.60125
SVM          0.59950
```

```
In [ ]: import matplotlib.pyplot as plt
import seaborn as sns

# Create a list of method names and their corresponding accuracies
methods = ['SVM (HOG)', 'CNN (CNN)']
accuracies = [accuracy_svm, accuracy_cnn]

# Create a bar plot to visualize the accuracies
plt.figure(figsize=(8, 6))
sns.barplot(x=methods, y=accuracies)
plt.title('Accuracy Comparison: SVM vs CNN')
plt.xlabel('Method')
plt.ylabel('Accuracy')
plt.ylim([0, 1]) # Set the y-axis limit from 0 to 1
plt.show()
```



Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js